

# CSGE602055 Operating Systems

## CSF2600505 Sistem Operasi

### Week 05: Virtual Memory

Rahmat M. Samik-Ibrahim

University of Indonesia

<https://os.vlsm.org/>

Always check for the latest revision!

REV163 13-Sep-2018

# Operating Systems 2018-2 (Room 3114)

## R/M (Tu/Th 13-15) | I (Tu/Th 15-17)

| Week           | Schedule                            | Topic                                       | OSC10                                  |
|----------------|-------------------------------------|---|--|
| Week 00        | 04 Sep - 12 Sep 2018                | Overview 1, Virtualization & Scripting      | Ch. 1, 2, 18.                          |
| Week 01        | 13 Sep - 19 Sep 2018                | Overview 2, Virtualization & Scripting      | Ch. 1, 2, 18.                          |
| Week 02        | 20 Sep - 26 Sep 2018                | Security, Protection, Privacy, & C-language | Ch. 16, 17                             |
| Week 03        | 27 Sep - 03 Oct 2018                | File System & FUSE                          | Ch. 13, 14, 15                         |
| Week 04        | 04 Oct - 10 Oct 2018                | Addressing, Shared Lib, & Pointer           | Ch. 9                                  |
| Week 05        | 11 Oct - 17 Oct 2018                | Virtual Memory                              | Ch. 10                                 |
| Reserved       | 18 Oct - 19 Oct 2018                |   |  |
| Mid-Term       | 20 Oct - 27 Oct 2018                | MidTerm (UTS): TBA                          |  |
| Week 06        | 30 Oct - 05 Nov 2018                | Concurrency: Processes & Threads            | Ch. 3, 4                               |
| Week 07        | 06 Nov - 12 Nov 2018                | Synchronization & Deadlock                  | Ch. 6, 7, 8                            |
| Week 08        | 13 Nov - 21 Nov 2018                | Scheduling                                  | Ch. 5                                  |
| Week 09        | 22 Nov - 28 Nov 2018                | Disks, BIOS, Loader, & Systemd              | Ch. 11                                 |
| Week 10        | 29 Nov - 05 Dec 2018                | I/O & Programming                           | Ch. 12                                 |
| Reserved       | 06 Dec - 14 Dec 2018                |   |  |
| Final<br>Extra | 15 Dec - 22 Dec 2018<br>12 Jan 2019 | Final (UAS): TBA<br>Extra assignment        | This schedule is<br>subject to change. |

# The Weekly Check List

- ☐ **Resources:** <https://os.vlsm.org/>
  - ☐ **(THIS) Slides** — <https://github.com/UI-FASILKOM-OS/SistemOperasi/tree/master/pdf/>
  - ☐ **Demos** — <https://github.com/UI-FASILKOM-OS/SistemOperasi/tree/master/demos/>
  - ☐ **Extra** — [BADAK.cs.ui.ac.id:///extra/](http://BADAK.cs.ui.ac.id:///extra/)
  - ☐ **Problems** — [rms46.vlsm.org/2/195.pdf](http://rms46.vlsm.org/2/195.pdf), [196.pdf](http://rms46.vlsm.org/2/196.pdf), ..., [205.pdf](http://rms46.vlsm.org/2/205.pdf)
- ☐ **Text Book:** any recent/decent OS book. Eg. **(OSC10)** Silberschatz et. al.: **Operating System Concepts**, 10<sup>th</sup> Edition, 2018.
- ☐ Encode your **QRC** with image size of approximately 250x250 pixels: **"OS182 CLASS ID SSO-ACCOUNT Your-Full-Name"**
- ☐ For **Week 00**, send your **embedded QRC before the 2<sup>nd</sup> lecture**  
<mailto:operatingsystems@vlsm.org>  
With Subject: OS182 CLASS ID SSO-ACCOUNT Your-Full-Name
- ☐ Write your Memo (with QRC) **every week**.
- ☐ Login to [badak.cs.ui.ac.id](http://badak.cs.ui.ac.id) via [kawung.cs.ui.ac.id](http://kawung.cs.ui.ac.id) for at least **10 minutes** every week. Copy the weekly demo files to your own home directory.  
Eg. (Week00): `cp -r /extra/Week00/W00-demos/ W00-demos/`

# Week 05: Memory

- 1 Start
- 2 Schedule
- 3 Week 05
- 4 Week 05
- 5 Virtual Memory
- 6 Memory Allocation Algorithm
- 7 TOP
- 8 06-memory
- 9 The End

# Week 05 Virtual Memory: Topics<sup>1</sup>

- Review of physical memory and memory management hardware
- Virtual Memory
- Caching
- Memory Allocation
- Memory Performance
- Working sets and thrashing

---

<sup>1</sup>Source: ACM IEEE CS Curricula 2013

# Week 05 Virtual Memory: Learning Outcomes<sup>1</sup>

- Explain memory hierarchy and cost-performance trade-offs. [Familiarity]
- Summarize the principles of virtual memory as applied to caching and paging. [Familiarity]
- Describe the reason for and use of cache memory (performance and proximity, different dimension of how caches complicate isolation and VM abstraction). [Familiarity]
- Defend the different ways of allocating memory to tasks, citing the relative merits of each. [Assessment]
- Evaluate the trade-offs in terms of memory size (main memory, cache memory, auxiliary memory) and processor speed. [Assessment]
- Discuss the concept of thrashing, both in terms of the reasons it occurs and the techniques used to recognize and manage the problem. [Familiarity]

---

<sup>1</sup>Source: ACM IEEE CS Curricula 2013

# Virtual Memory

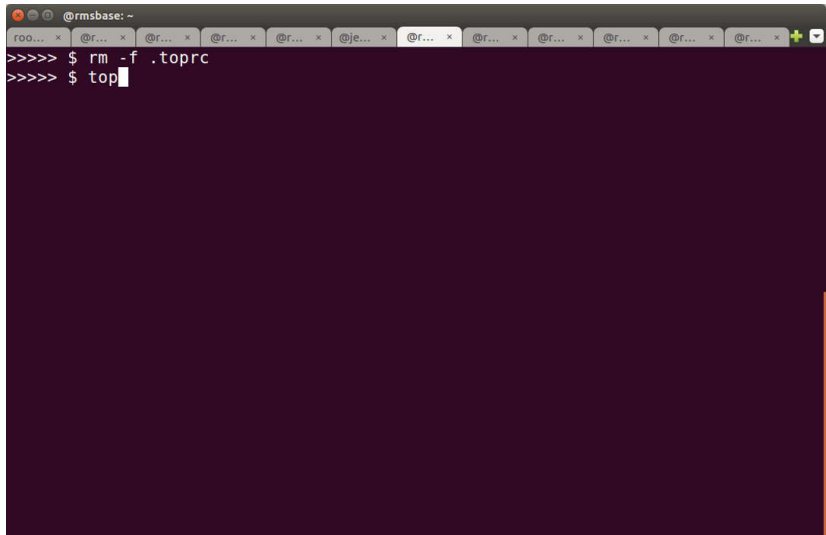
- Reference: (OSC10-ch10 demo-w05)
- Virtual Memory: Separation Logical from Physical.
- Virtual Address Space: logical view.
- Demand Paging
- Page Flags: Valid / Invalid
- Page Fault
- Demand Paging Performance
- Copy On Write (COW)
- Page Replacement Algorithm
  - Reference String
  - First-In-First-Out (FIFO)
  - Belady Anomaly
  - Optimal Algorithm
  - Least Recently Used (LRU)
  - LRU Implementation
  - Least Frequently Used (LFU)
  - Most Frequently Used (MFU)

# Allocation Algorithm

- Page-Buffering Algorithms
- Allocation of Frames
- Fixed Allocation
- Priority Allocation
- Global vs. Local Allocation
- Non-Uniform Memory Access (NUMA)
- Thrashing
- Working-Set Model
- Shared Memory via Memory-Mapped I/O
- Kernel
  - Buddy System Allocator
  - Slab Allocator



# TOP



A terminal window titled "@rmsbase: ~" with multiple tabs. The terminal shows the following commands and output:

```
>>>>> $ rm -f .toprc
>>>>> $ top
```

The terminal output is currently blank, indicating that the 'top' command has been executed but its output has not yet been displayed.

Figure: top

# TOP (2)

```
@rmsbase: ~
top - 18:37:28 up 14:07, 1 user, load average: 2.77, 2.71, 2.74
Tasks: 128 total, 1 running, 127 sleeping, 0 stopped, 0 zombie
%Cpu(s): 14.6 us, 17.2 sy, 0.0 ni, 68.1 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 8197060 total, 935152 used, 7261908 free, 191512 buffers
KiB Swap: 683004 total, 0 used, 683004 free. 639140 cached Mem
```

| PID  | USER | PR | NI  | VIRT   | RES  | SHR  | S | %CPU  | %MEM | TIME+    | COMMAND     |
|------|------|----|-----|--------|------|------|---|-------|------|----------|-------------|
| 518  | root | 20 | 0   | 162032 | 112  | 0    | S | 225.2 | 0.0  | 1882:33  | rngd        |
| 3448 | root | 20 | 0   | 0      | 0    | 0    | S | 14.0  | 0.0  | 0:09.14  | kworker/0:2 |
| 3198 | root | 20 | 0   | 0      | 0    | 0    | S | 9.6   | 0.0  | 5:29.03  | kworker/4:0 |
| 3062 | root | 20 | 0   | 0      | 0    | 0    | S | 5.0   | 0.0  | 11:55.39 | kworker/1:2 |
| 3289 | root | 20 | 0   | 0      | 0    | 0    | S | 2.3   | 0.0  | 3:41.00  | kworker/6:1 |
| 7    | root | 20 | 0   | 0      | 0    | 0    | S | 2.0   | 0.0  | 1:08.44  | rcu_sched   |
| 3376 | root | 20 | 0   | 0      | 0    | 0    | S | 1.3   | 0.0  | 0:18.73  | kworker/5:0 |
| 1914 | root | 20 | 0   | 0      | 0    | 0    | S | 0.3   | 0.0  | 13:10.69 | kworker/2:1 |
| 1    | root | 20 | 0   | 28684  | 4736 | 3012 | S | 0.0   | 0.1  | 0:02.91  | systemd     |
| 2    | root | 20 | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:00.01  | kthreadd    |
| 3    | root | 20 | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:15.26  | ksoftirqd/0 |
| 5    | root | 0  | -20 | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:00.00  | kworker/0:+ |
| 8    | root | 20 | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:00.00  | rcu_bh      |
| 9    | root | rt | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:00.00  | migration/0 |
| 10   | root | rt | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:00.25  | watchdog/0  |
| 11   | root | rt | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:00.28  | watchdog/1  |
| 12   | root | rt | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:00.00  | migration/1 |
| 13   | root | 20 | 0   | 0      | 0    | 0    | S | 0.0   | 0.0  | 0:06.80  | ksoftirqd/1 |

Figure: "h" = help

# TOP (3)

```
@rmsbase: ~
roo... x @f... x @f... x @f... x @f... x @je... x @f... x @f... x @f... x @f... x @f... x @f... x +
Fields Management for window 1:Def, whose current sort field is %CPU
Navigate with Up/Dn, Right selects for move then <Enter> or Left commits,
'd' or <Space> toggles display, 's' sets sort. Use 'q' or <Esc> to end!

* PID = Process Id      TTY = Controlling T      USED = Res+Swap Size
USER = Effective Use    TPGID = Tty Process G    nsIPC = IPC namespace
PR = Priority           SID = Session Id        nsMNT = MNT namespace
NI = Nice Value        nTH = Number of Thr     nsNET = NET namespace
VIRT = Virtual Image    P = Last Used Cpu       nsPID = PID namespace
RES = Resident Size     TIME = CPU Time         nsUSER = USER namespace
SHR = Shared Memory     SWAP = Swapped Size     nsUTS = UTS namespace
S = Process Statu      CODE = Code Size (Ki
%CPU = CPU Usage        DATA = Data+Stack (K
%MEM = Memory Usage     nMaj = Major Page Fa
TIME+ = CPU Time, hun   nMin = Minor Page Fa
COMMAND = Command Name/ nDRT = Dirty Pages C
PPID = Parent Proces    WCHAN = Sleeping in F
UID = Effective Use     Flags = Task Flags <s
RUID = Real User Id     CGROUPS = Control Group
RUSER = Real User Nam   SUPGIDS = Supp Groups I
SUID = Saved User Id    SUPGRPS = Supp Groups N
SUSER = Saved User Na   TGID = Thread Group
GID = Group Id          ENVIRON = Environment v
GROUP = Group Name      vMj = Major Faults
PRGP = Process Group    vMn = Minor Faults
```

Figure: Moving Fields: "f"

# TOP (4)

```
@rmsbase: ~
roo... x @f... x @f... x @f... x @f... x @je... x @f... x @f... x @f... x @f... x @f... x @f... x @f... x +
Fields Management for window 1:Def, whose current sort field is %CPU
Navigate with Up/Dn, Right selects for move then <Enter> or Left commits,
'd' or <Space> toggles display, 's' sets sort. Use 'q' or <Esc> to end!

* PID      = Process Id      SUID       = Saved User Id    vMn        = Minor Faults
* VIRT     = Virtual Image   SUSER      = Saved User Na    nsIPC      = IPC namespace
* RES      = Resident Size   GID        = Group Id        nsMNT      = MNT namespace
* SHR      = Shared Memory   GROUP      = Group Name      nsNET      = NET namespace
* SWAP     = Swapped Size    PGRP       = Process Group nsPID      = PID namespace
* CODE     = Code Size (Ki)  TTY        = Controlling T nsUSER     = USER namespac
* DATA    = Data+Stack (K)  TPGID      = Tty Process G nsUTS      = UTS namespace
* USED     = Res+Swap Size   SID        = Session Id
* nDRT     = Dirty Pages C   nTH        = Number of Thr
* PPID     = Parent Proces  P          = Last Used Cpu
%MEM       = Memory Usage   TIME       = CPU Time
USER       = Effective Use  nMaj       = Major Page Fa
PR         = Priority       nMin       = Minor Page Fa
NI         = Nice Value    WCHAN      = Sleeping in F
S          = Process Statu  Flags      = Task Flags <s
%CPU       = CPU Usage     CGROUPS    = Control Group
TIME+      = CPU Time, hun  SUPGIDS    = Supp Groups I
COMMAND    = Command Name/ SUPGRPS     = Supp Groups N
UID        = Effective Use  TGID       = Thread Group
RUID       = Real User Id   ENVIRON    = Environment v
RUSER      = Real User Nam vMj         = Major Faults
```

Figure: Moving Fields

# TOP (5)

```
@rmsbase: ~/Downloads
top - 19:57:14 up 11:38, 1 user, load average: 0.43, 0.54, 0.58
Tasks: 285 total, 2 running, 283 sleeping, 0 stopped, 0 zombie
%Cpu(s): 3.8 us, 1.3 sy, 0.0 ni, 94.6 id, 0.3 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 16385976 total, 269672 free, 3179788 used, 12936516 buff/cache
KiB Swap: 1000444 total, 994752 free, 5692 used. 12649780 avail Mem
```

| PID   | VIRT    | RES    | SHR    | SWAP | CODE   | DATA    | USED   | nDRT |
|-------|---------|--------|--------|------|--------|---------|--------|------|
| 3547  | 2377296 | 394828 | 165776 | 0    | 196    | 1642748 | 394828 | 0    |
| 1234  | 278216  | 87880  | 59116  | 0    | 2288   | 25164   | 87880  | 0    |
| 3321  | 2683572 | 433176 | 149376 | 0    | 196    | 1856708 | 433176 | 0    |
| 2708  | 1687448 | 214112 | 80608  | 0    | 12     | 1179008 | 214112 | 0    |
| 2841  | 679488  | 50860  | 30484  | 0    | 292    | 389096  | 50860  | 0    |
| 3748  | 1896812 | 321288 | 76656  | 0    | 133688 | 1474084 | 321288 | 0    |
| 3971  | 2047252 | 440112 | 97384  | 0    | 133688 | 1587052 | 440112 | 0    |
| 32501 | 630768  | 33500  | 27960  | 0    | 76     | 373220  | 33500  | 0    |
| 4067  | 8554396 | 320516 | 109756 | 0    | 196    | 7954584 | 320516 | 0    |
| 4130  | 2391592 | 341632 | 117636 | 0    | 196    | 1717824 | 341632 | 0    |
| 22635 | 2198448 | 274812 | 108000 | 0    | 196    | 1532152 | 274812 | 0    |
| 1292  | 0       | 0      | 0      | 0    | 0      | 0       | 0      | 0    |
| 2514  | 930224  | 34304  | 26028  | 0    | 36     | 448864  | 34304  | 0    |
| 3233  | 4515228 | 360812 | 126784 | 0    | 133688 | 3757984 | 360812 | 0    |
| 32495 | 33488   | 3380   | 2836   | 0    | 96     | 1264    | 3380   | 0    |
| 2388  | 44036   | 4424   | 2724   | 0    | 212    | 1716    | 4424   | 0    |
| 2412  | 423204  | 11380  | 5264   | 0    | 152    | 374232  | 11380  | 0    |
| 2512  | 685824  | 74188  | 36868  | 0    | 552    | 399836  | 74188  | 0    |

Figure: Write Configuration .toprc: "W"

# 06-memory

```
/* Copyright (C) 2016-2018 Rahmat M. Samik-Ibrahim
 * https://rahmatm.samik-ibrahim.vlsm.org/
 * This program is free script/software. This program is distributed in the
 * hope that it will be useful, but WITHOUT ANY WARRANTY; without even the
 * implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
 * REV04 Mon Mar 12 17:33:30 WIB 2018
 * START Mon Oct 3 09:26:51 WIB 2016
 */
#define MSIZE0 0x10000
#define MSIZE1 0x10008
#define MSIZE2 0x10009
#define MSIZE3 0x1000A
#define MSIZE4 0x20978
#define MSIZE5 0x20979
#define MSIZE6 0x2097A
#define MSIZE7 0xF0000
#define MSIZE8 0x10000
#define MSIZE9 0x1000
#define LINE 75
#define MAXSTR 80
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/types.h>

void printLine(int line) {
    while(line-- > 0) putchar('x');
    putchar('\n');
    fflush(NULL);
}
```

## 06-memory (2)

```
void main (void) {
    int  msize[] = {MSIZE0, MSIZE1, MSIZE2, MSIZE3, MSIZE4,
                    MSIZE5, MSIZE6, MSIZE7, MSIZE8, MSIZE9};

    int  ii, jj;
    int  myPID   = (int) getpid();
    char strSYS1[MAXSTR], strOUT[MAXSTR];
    char* chrStr  = strSYS1;
    char* chrPTR;

    printLine(LINE);
    sprintf(strSYS1, "top -b -n 1 -p%d | tail -5", myPID);
    system (strSYS1);
    sprintf(strSYS1, "top -b -n 1 -p%d | tail -1", myPID);
    for (ii=0; ii< (sizeof(msize)/sizeof(int)); ii++){
        chrStr = malloc(msize[ii]);
        fgets(strOUT, sizeof(strOUT)-1, popen(strSYS1, "r"));
        strOUT[(int) strlen(strOUT)-1]='\0';
        printf("%s [%X]\n", strOUT, msize[ii]);
        free(chrStr);
    }
    for (ii=0; ii< (sizeof(msize)/sizeof(int)); ii++){
        chrPTR = chrStr = malloc(msize[ii]);
        for (jj=0;jj<msize[ii];jj++)
            *chrPTR++='x';
        fgets(strOUT, sizeof(strOUT)-1, popen(strSYS1, "r"));
        strOUT[(int) strlen(strOUT)-1]='\0';
        printf("%s [%X]\n", strOUT, msize[ii]);
        free(chrStr);
    }
}
```

## 06-memory (2)

```
>>>>> $ ./06-memory
```

[illegible]

```
KiB Mem:  8197060 total,  957928 used,  7239132 free,  192520 buffers
```

```
KiB Swap: 683004 total, 0 used, 683004 free. 660108 cached
```

Mem

| PID  | VIRT | RES  | SHR  | SWAP | CODE | DATA | USED | nDRT      |
|------|------|------|------|------|------|------|------|-----------|
| 4362 | 4172 | 640  | 564  | 0    | 4    | 320  | 640  | 0         |
| 4362 | 4172 | 640  | 564  | 0    | 4    | 320  | 640  | 0 [10000] |
| 4362 | 4172 | 640  | 564  | 0    | 4    | 320  | 640  | 0 [10008] |
| 4362 | 4308 | 640  | 564  | 0    | 4    | 456  | 640  | 0 [10009] |
| 4362 | 4244 | 1176 | 1068 | 0    | 4    | 392  | 1176 | 0 [1000A] |
| 4362 | 4244 | 1176 | 1068 | 0    | 4    | 392  | 1176 | 0 [20978] |
| 4362 | 4376 | 1176 | 1068 | 0    | 4    | 524  | 1176 | 0 [20979] |
| 4362 | 4376 | 1192 | 1068 | 0    | 4    | 524  | 1192 | 0 [2097A] |
| 4362 | 5340 | 1192 | 1068 | 0    | 4    | 1488 | 1192 | 0 [F0000] |
| 4362 | 4376 | 1200 | 1068 | 0    | 4    | 524  | 1200 | 0 [10000] |
| 4362 | 4376 | 1200 | 1068 | 0    | 4    | 524  | 1200 | 0 [1000]  |



## 06-memory (3)

|      |      |      |      |   |   |      |      |           |
|------|------|------|------|---|---|------|------|-----------|
| 4362 | 4376 | 1200 | 1068 | 0 | 4 | 524  | 1200 | 0 [1000]  |
| 4362 | 4376 | 1200 | 1068 | 0 | 4 | 524  | 1200 | 0 [10000] |
| 4362 | 4376 | 1276 | 1068 | 0 | 4 | 524  | 1276 | 0 [10008] |
| 4362 | 4376 | 1276 | 1068 | 0 | 4 | 524  | 1276 | 0 [10009] |
| 4362 | 4376 | 1284 | 1068 | 0 | 4 | 524  | 1284 | 0 [1000A] |
| 4362 | 4376 | 1284 | 1068 | 0 | 4 | 524  | 1284 | 0 [20978] |
| 4362 | 4376 | 1352 | 1068 | 0 | 4 | 524  | 1352 | 0 [20979] |
| 4362 | 4376 | 1352 | 1068 | 0 | 4 | 524  | 1352 | 0 [2097A] |
| 4362 | 5340 | 2144 | 1068 | 0 | 4 | 1488 | 2144 | 0 [F0000] |
| 4362 | 5340 | 2324 | 1068 | 0 | 4 | 1488 | 2324 | 0 [10000] |
| 4362 | 5340 | 2324 | 1068 | 0 | 4 | 1488 | 2324 | 0 [1000]  |

>>>>> \$

# The End

- ☐ This is the end of the presentation.
- ☒ This is the end of the presentation.
  - This is the end of the presentation.